

Saturated Superfluid Vacuum I

Topological Defects in a Saturated Superfluid Vacuum:
The Geometric Origin of Matter and
the α -Harmonic Mass Ladder

Stig Norland
Independent Researcher
Bergen, April 2026

Abstract

We propose that the physical vacuum is a compressible, inviscid superfluid near saturation density — a single postulate from which the structure of fundamental particles and constraints on the fine-structure constant emerge. Particles are stable topological defects: quantised toroidal vortex rings (leptons) and spherical breathers stabilised by trefoil vortex skeletons (baryons).

From the logarithmic Schrödinger Lagrangian we derive three results from first principles. First, the sound-speed condition $c_s = c$ fixes the stiffness parameter $b = m_0 c^2 / (2\rho_0)$ and the healing length $\xi = \hbar / (m_0 c)$. Second, minimisation of the three-term vortex-ring energy functional over torus radius R and core radius a yields the unique equilibrium $R_e^* = \xi / \alpha$, from which the vacuum density ρ_0 is fixed via the Lamb formula with no free parameters beyond the two empirical inputs m_e and α . Third, the ratio of the chiral-shear speed $c_\perp$ to the sound speed c_s identifies $\alpha \approx 1/137$ as a geometric property of the medium.

These results establish the base energy scale $\mu_0 \equiv m_e c^2 / \alpha \approx 70$ MeV as the energy of a minimal vortex filament segment of the electron Compton circumference. The charged pion, as the minimal vortex-antivortex bound state carrying two winding numbers, then has mass $m_{\pi^\pm} = 2\mu_0$, agreeing with CODATA at 0.34%.

Direct 3D toroidal Bogoliubov–de Gennes projection with chiral-shear coupling, developed in this paper, identifies the dynamical mechanism behind the muon rung: the lowest hybrid of the $m_\varphi = 0$ core-breathing mode with the $m_\varphi = \pm 1$ Kelvin centerline sector in the helicity basis, coupled through the second variation of $\mathcal{L}_\perp$. Driven-response calculations confirm a selection rule: core-breathing boundary drives preferentially excite the $(3/2)\mu_0$ tongue identified with the muon, while Kelvin helicity drives preferentially excite the $2\mu_0$ tongue identified with the charged pion. In the present meridional projection the muon eigenfrequency falls in the range $\omega_\mu / \omega_c \in [0.19, 0.23]$ depending on meridional domain size, bracketing the observed value $0.207\omega_c$ without yet resolving to a single number; thin-ring analytic closure is required for a sharp prediction and is deferred to Paper II. For the proton, the empirical scaling $m_p c^2 = N_Y \cdot F \cdot \mu_0$ with $N_Y \approx 3.007$ and form factor $F \approx 4.4$ – 4.47 at the inter-strand half-spacing cutoff $R \approx 1.18\xi$ yields $m_p \approx 927$ – 941 MeV (1.2%–0.3% from CODATA, depending on the cutoff convention).

Two empirical inputs (m_e, α) are used. The remaining predictions are consistent with the observed lepton and light-hadron spectrum at the $\lesssim 1.5\%$ level under the present effective theory, with explicit dependence on the regularisation cutoff R and the meridional domain size L_{box} documented in the body. A dedicated Claim Status section (§8) tags every quantitative claim as *derived*, *structural*, *candidate*, or *ansatz*, and links each candidate-grade claim to the specific GitHub issue (#13, #16) whose computation would change its status. The open problems for Paper II are: thin-ring analytic closure of the muon eigenfrequency, full 3D minimisation for the proton form factor (including a cutoff-stability scan for $F(R)$ and a dedicated N_Y cross-grid refinement), and a basis-enrichment programme that demonstrates the $1/\sqrt{N_{\text{modes}}}$ residual scaling proposed for the muon rung. The previously headlined 0.3% proton-mass agreement is the precision at a specific calibration cutoff $R \approx 1.18\xi$ and is reframed in this paper as a $\sim 1.5\%$ match with documented cutoff sensitivity.

Contents

1	Introduction: The Crisis of Abstraction	3
2	The Fundamental Postulate	3
2.1	Effective Lagrangian and Equations of Motion	4
3	The Geometric Origin of Matter	5
3.1	The Electron (The Torus)	6
3.2	The Proton (The Sphere)	8
3.3	The Neutron — A Speculative Extension	9
4	The α-Harmonic Mass Ladder	9
4.1	The Empirical Ladder	9
4.2	Physical origin of the base scale	10
4.3	Status of individual rungs	10
4.4	The Charged Pion ($\pi^\pm$)	10
4.5	The Muon ($\mu^\pm$)	11
4.6	Higher rungs and the proton	12
5	Forces: Hydrodynamic Emergence (Preview)	13
6	Outlook and Limitations	13
7	Conclusion and Next Steps	15
8	Claim Status	16
9	Acknowledgments	18
A	Numerical Construction of the Trefoil Y-Junction Breather	19
B	Numerical Minimisation of the Trefoil Y-Junction Energy	19
C	BdG Eigenvalue Problem: Equations and Definitive Numerical Result	21
D	Fat-Torus Self-Interaction Coefficient	23
E	Helicity-Basis Kelvin Bridge and Driven-Response Selection Rule	23
E.1	Projection basis	23
E.2	Current-curl second-variation bridge	24
E.3	Driven-response selection rule	24
E.4	Box-size dependence and residual uncertainty	25
E.5	Interpretation and comparison to experimental precision	25

1 Introduction: The Crisis of Abstraction

The Standard Model predicts particle masses, charges and interaction rates with extraordinary precision, yet it provides no mechanism for *why* these quantities take their observed values. Some 19 free parameters—particle masses, coupling constants, mixing angles—must be inserted by hand. At the same time, attempts to quantize gravity perturbatively produce non-renormalizable divergences, and roughly 95% of the universe’s energy content must be assigned to unknown dark components inferred only from gravitational anomalies.

We propose that this impasse reflects a missing ontological layer. The vacuum is not empty: it is a physical substance—a saturated superfluid. All phenomena emerge from its hydrodynamics. This idea has deep roots. Volovik has shown that the low-energy physics of superfluid ^3He reproduces the structure of the Standard Model, with Lorentz symmetry, gauge bosons, and chiral fermions emerging as collective excitations [1]. Barceló, Liberati, and Visser established the analogue-gravity programme: acoustic perturbations of any irrotational, barotropic fluid obey a curved-spacetime wave equation, so gravitational phenomenology emerges from fluid dynamics [3]. Hu has argued that spacetime itself is an emergent concept from a more fundamental substrate [4].

The present work pursues a specific realisation of this programme: a logarithmic nonlinear Schrödinger (LogSE) field theory for the vacuum order parameter, supplemented by a chiral-shear term. Stable topological defects of this field—toroidal vortex rings and spherical breathers—are identified with fundamental particles. The mass spectrum of these defects organises into a harmonic ladder governed by the fine-structure constant $\alpha \approx 1/137$.

This paper is organised as follows. Section 2 states the fundamental postulate, derives the equations of motion from the LogSE Lagrangian, derives the sound speed, and establishes the origin of α from the chiral-shear coupling. Section 3 establishes the equilibrium vortex geometry and the vacuum density from the Lamb vortex-ring formula, taking m_e and α as empirical inputs, and presents the proton and neutron models. Section 4 develops the α -harmonic mass ladder, including the pion (analytically), the muon (dynamically motivated bracket), and the proton (numerically extracted, cutoff-dependent). Section 5 sketches the emergence of forces. Sections 6–7 give the outlook and conclusions. Appendices contain the numerical minimisation and the BdG eigenvalue equations together with their numerical solution.

2 The Fundamental Postulate

The physical vacuum is a continuous, compressible, inviscid superfluid at near-saturation density ρ_0 , with background pressure $P_0 \gg \rho_0 c^2$. All known physics emerges from the hydrodynamics of this single medium; no additional fields or symmetry-breaking mechanisms are required at the foundational level.

The defining characteristics are:

- **Compressibility and high stiffness.** The equation of state is such that small density perturbations propagate as pressure waves at speed $c_s = \sqrt{\partial P / \partial \rho}$. At the background state, we require $c_s = c$ (the observed speed of light) to be constant (Lorentz-invariant in the low-momentum limit). The condition $P_0 \gg \rho_0 c^2$ ensures the medium remains far from the laminar-to-turbulent transition under ordinary stresses, while still allowing topological defects to form under extreme local strain.
- **Inviscid flow.** Zero (or exponentially suppressed) viscosity guarantees dissipation-free long-range coherence—persistent currents, quantized circulation, and stable solitons. This is the hydrodynamic analogue of unitarity and conservation laws.

- **Light as sound.** Electromagnetic waves (and massless modes generally) are long-wavelength longitudinal/transverse pressure oscillations in the superfluid. Thus,

$$c = c_s(\rho_0) = \sqrt{\left. \frac{\partial P}{\partial \rho} \right|_{\rho_0}}. \quad (1)$$

In logarithmic superfluid models (common in SVT literature), the pressure-density relation takes the form $P \propto -\rho \ln(\rho/\bar{\rho})$ [1, 2], yielding constant c_s at low excitations while allowing deformed dispersion at high momenta.

- **Special relativity as emergent hydrodynamics.** Objects composed of stable topological defects (vortices/breathers) experience length contraction and time dilation as equilibrium responses to motion through the medium. A moving defect induces a wake; to minimise energy/disruption the defect contracts in the propagation direction (Lorentz–FitzGerald style). The medium itself remains undetectable in Michelson–Morley-type interferometers because the “ether wind” is canceled by the defect’s own flow field—the drag is internal to the object-medium system.

This postulate recovers special relativity in the low-velocity, low-excitation limit without assuming a priori Lorentz symmetry [1, 3]; the symmetry emerges from the constancy of c_s and the irrotational nature of small perturbations (potential flow). At higher energies/momenta, deviations from exact Lorentz invariance are expected (deformed dispersion relations), consistent with hints from ultra-high-energy cosmic rays and some quantum-gravity phenomenology [3, 10].

2.1 Effective Lagrangian and Equations of Motion

The dynamics of the saturated superfluid vacuum are governed by an effective field theory for the complex order-parameter field Ψ (normalised so $|\Psi|^2 = \rho/\rho_0$). Following the superfluid-vacuum framework of Volovik [1], we adopt a logarithmic nonlinear Schrödinger Lagrangian:

$$\mathcal{L} = \frac{i\hbar}{2}(\Psi^* \partial_t \Psi - \Psi \partial_t \Psi^*) - \frac{\hbar^2}{2m_0} |\nabla \Psi|^2 - V(\rho), \quad (2)$$

where the potential is

$$V(\rho) = -b\rho [\ln(\rho/\bar{\rho}) - 1] + V_0, \quad (3)$$

with $b > 0$ a stiffness parameter and $\bar{\rho}$ the saturation density. Applying the Euler–Lagrange equation to (2), using $\partial V/\partial \Psi^* = (dV/d\rho)\rho_0\Psi$ and $dV/d\rho = -b\ln(\rho/\bar{\rho})$, yields the logarithmic Schrödinger equation (LogSE):

$$i\hbar \partial_t \Psi = -\frac{\hbar^2}{2m_0} \nabla^2 \Psi - b\rho_0 \ln\left(\frac{\rho_0 |\Psi|^2}{\bar{\rho}}\right) \Psi. \quad (4)$$

Madelung representation. Writing $\Psi = \sqrt{\rho/\rho_0} e^{i\theta}$ with superfluid velocity $\mathbf{v} = (\hbar/m_0)\nabla\theta$, the LogSE (4) separates into real and imaginary parts. The imaginary part gives the exact continuity equation

$$\partial_t \rho + \nabla \cdot (\rho \mathbf{v}) = 0. \quad (5)$$

The real part, after dividing by ρ , gives an Euler equation

$$\partial_t \mathbf{v} + (\mathbf{v} \cdot \nabla) \mathbf{v} = -\frac{1}{m_0 \rho} \nabla P(\rho) + \frac{\hbar^2}{2m_0^2} \nabla \left(\frac{\nabla^2 \sqrt{\rho}}{\sqrt{\rho}} \right), \quad (6)$$

where the pressure follows the logarithmic equation of state $P(\rho) = b\rho \ln(\rho/\bar{\rho})$. The last term is the quantum-pressure (Bohm potential), which is negligible at wavelengths $\lambda \gg \xi$; at those

scales the LogSE reduces to a *compressible inviscid fluid* with logarithmic EOS. The thermodynamic sound speed $c_{\text{thermo}} = \sqrt{(1/m_0)dP/d\rho|_{\rho_0}} = \sqrt{b/m_0}$ is half the Bogoliubov result $c_s = \sqrt{2b\rho_0/m_0}$; the discrepancy is resolved by the coupled amplitude-phase Bogoliubov dispersion, which gives $\omega^2 = c_s^2 k^2(1 + k^2 \xi^2)$ with $c_s = \sqrt{2b\rho_0/m_0}$.

Sound speed. Linearising around the uniform background $\Psi = 1 + \eta$, $|\eta| \ll 1$, with $\rho_0 = \bar{\rho}$, the density channel satisfies a wave equation with phase velocity

$$c_s = \sqrt{\frac{2b\rho_0}{m_0}}. \quad (7)$$

Setting $c_s = c$ fixes the stiffness parameter: $b = m_0 c^2/(2\rho_0)$. The healing length (vortex core size) is

$$\xi = \frac{\hbar}{\sqrt{2m_0 b \rho_0}} = \frac{\hbar}{m_0 c}, \quad (8)$$

i.e., the Compton wavelength of the reference mass m_0 .

Transverse stiffness and the origin of α . The LogSE admits only irrotational flow at linear order. To support stable vortices and a transverse shear channel (necessary for electromagnetism to emerge as chiral flow gradients), we supplement $\mathcal{L}$ with the chiral-shear term

$$\mathcal{L}_\perp = -\frac{\lambda \hbar^2}{2m_0 \rho_0^2} |(\nabla \times \mathbf{j})|^2, \quad (9)$$

where $\mathbf{j} = (\hbar/2m_0 i)(\Psi^* \nabla \Psi - \Psi \nabla \Psi^*)$ is the probability current and $\lambda > 0$ is a dimensionless coupling. This produces a transverse mode with phase speed $c_\perp(k) \propto k$ at wavenumber k . Evaluating at the vortex core scale $k \sim 1/\xi$ gives

$$\frac{c_\perp}{c} = \sqrt{\frac{\lambda m_0}{\rho_0}} \equiv \alpha. \quad (10)$$

The fine-structure constant $\alpha \approx 1/137$ is therefore the ratio of the transverse shear speed to the longitudinal sound speed in the vacuum at the core scale. Equation (10) fixes the chiral coupling $\lambda = \alpha^2 \rho_0/m_0$; with $\alpha = 1/137$ and $\lambda \approx 2000$ (from the numerical minimisation in Appendix A), this yields $\rho_0/m_0 = \lambda/\alpha^2 \approx 3.75 \times 10^7$ as a dimensionless characterisation of the vacuum state.

Two derived results

1. The LogSE with potential (3) has longitudinal sound speed $c_s = \sqrt{2b\rho_0/m_0}$; setting $c_s = c$ gives $b = m_0 c^2/(2\rho_0)$ and healing length $\xi = \hbar/(m_0 c)$.
2. With the chiral-shear term (9), the transverse phase speed at the core scale is $c_\perp = c\sqrt{\lambda m_0/\rho_0}$; identifying $c_\perp/c = \alpha$ fixes $\lambda = \alpha^2 \rho_0/m_0$.

These are the only two results needed to motivate $\mu_0 = m_e c^2/\alpha$ as the natural flux-tube energy scale (see §4).

This Lagrangian admits **quantised circulation** $\oint \mathbf{v} \cdot d\mathbf{l} = nh/m_0$ (Onsager–Feynman) and supports stable topological defects. Stable defect geometries are constructed and their mass assignments established in Sections 3–4.

3 The Geometric Origin of Matter

Matter forms when the medium is stressed beyond its laminar limit, collapsing into stable topological defects (vortices).

3.1 The Electron (The Torus)

The electron is modelled as a quantized toroidal vortex ring in the saturated superfluid vacuum—a stable, self-sustaining topological defect with finite size.

- **Spin.** The intrinsic angular momentum arises from the quantized circulation around the toroidal path:

$$\oint \mathbf{v} \cdot d\mathbf{l} = \kappa_0 = \frac{h}{m_0} = n \frac{h}{m_e}, \quad n = 1, \quad (11)$$

where κ_0 is the quantum of circulation (Onsager–Feynman quantization), h is Planck’s constant, and m_0 is a *defect-relative reference mass*: it takes the rest mass of the defect species under consideration. For the electron (toroidal vortex ring) $m_0 = m_e$; for the proton (knotted breather, analysed in Paper II) $m_0 = m_p$. No single universal value is postulated for the medium as a whole.

Each defect carries its own healing length $\xi \equiv \hbar/(m_0 c)$: the electron’s is $\xi_e = \hbar/(m_e c) \approx 3.86 \times 10^{-13}$ m (the reduced Compton wavelength $\bar{\lambda}_e$), and the proton’s is $a_p \equiv \hbar/(m_p c) \approx 2.10 \times 10^{-16}$ m. The ratio $a_p/\xi_e = m_e/m_p \approx 1/1836$ is a *cross-defect* quantity; it governs the proton breather’s sub-grain coupling to the long-range phonon field in Paper II’s gravity section. The medium’s domain of validity as a continuum effective theory extends to $\ell_{\text{grain}} \equiv a_p$; below this scale the LogSE description is silent, in the same sense that Navier–Stokes is silent about sub-molecular flow (see Paper II §2). Spin $s = \hbar/2$ follows when the circulation quantum is combined with the half-quantum topology of the ring; the detailed Finkelstein–Rubinstein argument (π_1 of the configuration space is $\mathbb{Z}_2$, placing the torus in the fermionic sector) will be given in Paper II.

- **Charge.** Electric charge emerges from the *chirality* (handedness) of the toroidal flow. Right-handed circulation corresponds to positive charge; the opposite handedness yields the positron. The magnitude $|e|$ is tied to the transverse flow field strength at large distances, which behaves as a dipole-like circulation decaying as $1/r^2$, reproducing the Coulomb law in the far field.
- **Equilibrium radius (analytic derivation, two structural inputs).** The electron radius R_e is determined by minimising the total hydrodynamic energy of the vortex ring over the two geometric parameters (R, a) , where R is the major (centreline) radius and a is the core radius. Three physical contributions are included. The minimisation itself is analytic (status: *derived*; §8), but the chiral-shear coefficient appearing below (Eq. 13) is a structural input fixed by the $c_\perp/c = \alpha$ identification of §2, not a separately computed quantity.

(i) *Kinetic energy (Lamb formula).* For a thin ring ($a \ll R$) carrying one quantum of circulation $\kappa_0 = h/m_0$:

$$E_{\text{kin}}(R, a) = \frac{1}{2} \rho_0 \kappa_0^2 R \left(\ln \frac{8R}{a} - \frac{7}{4} \right). \quad (12)$$

(ii) *Core self-energy (GPE).* Minimising over a alone gives $a^* = \xi = \hbar/(m_0 c)$ with no free parameters.

(iii) *Chiral-shear energy (LogSE).* The chiral-shear term (9) with stiffness $\lambda = \alpha^2 \rho_0 / m_0$ contributes an energy proportional to the ring circumference that resists expansion, with coefficient satisfying

$$2\pi\gamma\alpha^2 = \Lambda + 1 = \ln\left(\frac{8}{\alpha}\right) - \frac{3}{4} \approx 6.25, \quad (13)$$

where $\Lambda \equiv \ln(8/\alpha) - 7/4 \approx 5.25$.

Total functional and minimisation. Setting $a = a^* = \xi$ and working in units of $\frac{1}{2}\rho_0\kappa_0^2\xi$ with $r \equiv R/\xi$:

$$\mathcal{E}(r) = r \left(\ln 8r - \frac{7}{4} \right) + \frac{2C}{r} - (\Lambda + 1) \alpha^2 r, \quad (14)$$

where $C = 1/8$ from the elliptic self-interaction integral (Appendix D). Setting $d\mathcal{E}/dr = 0$:

$$\ln 8r^* - \frac{3}{4} - \frac{2C}{r^{*2}} = (\Lambda + 1) \alpha^2. \quad (15)$$

At $r^* \sim 1/\alpha \gg 1$ the term $2C/r^{*2} \sim \alpha^2$ is negligible. Using $(\Lambda + 1)\alpha^2 = \ln(8/\alpha) - \frac{3}{4}$, Eq. (15) reduces to $\ln 8r^* = \ln(8/\alpha)$, with the unique solution

$$R_e^* = \frac{\xi}{\alpha} = \frac{\hbar}{\alpha m_0 c}. \quad (16)$$

This is the classical electron radius, derived without empirical input for R_e . Numerically, $d^2\mathcal{E}/dr^2|_{r^*} = \alpha > 0$: a *true minimum*, not a saddle.

Inserting R_e^* and $\Lambda \approx 5.25$ into (12):

$$m_e c^2 = \frac{1}{2} \rho_0 \kappa_0^2 \frac{\xi}{\alpha} \Lambda, \quad (17)$$

which fixes the vacuum saturation density:

$$\rho_0 = \frac{2\alpha\Lambda m_e^4 c^3}{\pi^2 \hbar^3} \approx 1.9 \frac{m_e^4 c^3}{\hbar^3}. \quad (18)$$

The derivation uses two empirical inputs (m_e , α) and no free parameters beyond those already fixed by the Lagrangian.

Equilibrium radius: logical chain

1. GPE fixes $a^* = \xi$ (no thin-core assumption).
2. Chiral stiffness $\lambda = \alpha^2 \rho_0 / m_0$ adds circumference energy $\propto \alpha^2 R$ resisting expansion.
3. Balance of kinetic ($\propto R \ln R$) against chiral ($\propto \alpha^2 R$) yields a unique minimum at $R_e^* = \xi/\alpha$.
4. α enters once, from the Lagrangian; R_e^* is not assumed.

- **Zitterbewegung.** The observed rapid trembling motion of the Dirac electron (frequency $\sim 2m_e c^2/\hbar$) is reinterpreted mechanically as the vortex ring *surfing* its own Kelvin waves—helical perturbations that propagate along the toroidal core. The Kelvin-wave dispersion relation on a quantized vortex line is

$$\omega(k) \approx \frac{\kappa_0 k^2}{4\pi} \left(\ln \frac{1}{ka} + \gamma \right), \quad (19)$$

where k is wavenumber and γ a constant. For short-wavelength modes ($ka \sim 1$), the frequency approaches the Compton scale, producing the characteristic jitter without invoking negative-energy seas or pair creation.

This toroidal defect is topologically stable: small perturbations radiate Kelvin waves but the core persists.

Gyromagnetic ratio. The electron vortex ring carries angular momentum S from two sources: the topological winding number ($S_{\text{topo}} = \hbar/2$, from the fermionic sector of the configuration space, cf. Paper II) and the orbital flow (L_{orb}). The magnetic moment of a current loop is $\mu = I \cdot \pi R_e^2$, where the effective current is $I = (e/h)\kappa_0$ (one unit of charge completing one circulation quantum per unit time). Using $\kappa_0 = h/m_e$:

$$\mu = \frac{e}{h} \cdot \frac{h}{m_e} \cdot \pi R_e^2 = \frac{e\pi R_e^2}{m_e}. \quad (20)$$

The total angular momentum associated with the flow is $L = \rho_0 \kappa_0 \pi R_e^2$ (standard vortex-ring result), but the physical spin controlling the gyromagnetic ratio is $S_{\text{topo}} = \hbar/2$. The ratio

$$g \equiv \frac{2m_e c}{e} \cdot \frac{\mu}{S_{\text{topo}}} = \frac{2m_e c}{e} \cdot \frac{e\pi R_e^2/m_e}{\hbar/2} = \frac{4\pi c R_e^2}{\hbar} \quad (21)$$

This is parametrically of order $\xi \sim \hbar/(m_e c)$. The leading-order result $g \approx 2$ holds whenever $R_e \sim \xi$; higher-order Kelvin-wave dressing produces corrections $g - 2 \sim \alpha/\pi$, consistent with the observed anomalous magnetic moment. A precise derivation of $g = 2$ from the exact vortex normalisation is deferred to Paper II.

3.2 The Proton (The Sphere)

The proton is a spherical breather mode—a stable, oscillating cavity in the superfluid vacuum, prevented from collapse by an internal trefoil vortex skeleton.

The simplest three-dimensional structure that can stabilize a spherical void against the immense background pressure P_0 is a Y-junction [9] of three quantized vortex filaments meeting at a single central node (topologically a trefoil knot projected into three orthogonal directions). This configuration is the hydrodynamic origin of baryon number and confinement.

- **Quarks.** The three vortex filaments themselves. They are not independent particles but structural legs of the skeleton. Each carries a fraction of the total topological charge; their length and curvature set the proton radius (~ 0.84 fm from charge radius data).
- **Color charge.** The phase orientation of the three filaments at the central Y-junction. In three spatial dimensions the only stable balance is 120° separation in phase space. This reproduces the three “colors” and the requirement that the total wavefunction be color-neutral at the node. Reconnection or imbalance at the junction costs extra energy—the source of confinement.
- **Mass.** The rest energy is obtained by minimizing the total hydrodynamic energy functional (vortex-line tension + junction self-energy + acoustic breather energy):

$$m_p c^2 = N_Y \cdot \frac{m_e}{\alpha} \cdot F(\text{geometry}), \quad (22)$$

where $\mu_0 \equiv m_e/\alpha \approx 70$ MeV is the flux-tube mass scale and $N_Y \approx 3.007$ is the dimensionless node-cost factor from 3D numerical energy minimization of the Y-junction (Appendix B).

The form factor F encodes the ratio of the total 3D breather energy to the Y-junction line energy $N_Y \mu_0$. It can be written explicitly as

$$F = \frac{E_{\text{breather}}^{\text{3D}}}{N_Y \mu_0} = \frac{1}{N_Y \mu_0} \int \left[\frac{\hbar^2}{2m_0} |\nabla \Psi|^2 + V(|\Psi|^2) \right] d^3x, \quad (23)$$

where the integral is over the full 3D breather profile. From the 3D numerical minimisation with the trefoil Y-junction skeleton, $F \approx 4.47$. The consistency check is $N_Y \times F = 3.007 \times 4.47 \approx 13.44$, giving $m_p c^2 \approx 13.44 \times 70 \text{ MeV} \approx 941 \text{ MeV}$, within 0.3% of the observed value before fine relaxation of the cavity geometry. The full derivation of F from the 3D profile without any fitting will be given in Paper II.

Numerical update (2026-05-19, see Appendix B): A topology-preserving relaxation of Eq. (23) on the (2, 3)-trefoil knot has now been completed at three grid resolutions ($n^3 = 24^3, 48^3, 72^3$ with ξ -normalised box half-width 6ξ). With a grid-invariant straight-vortex calibration, F depends on the physical cutoff R at which the calibration tube terminates: $F(R=1.18\xi) \approx 4.4$ (fine-grid extrapolation) matches the analytic estimate; $R=1.18\xi$ is the natural inter-strand half-spacing of the trefoil with minor radius 0.85ξ . The product $N_Y \cdot F$ then gives $m_p \approx 927 \text{ MeV}$ (1.2% below the observed value), rather than the original 0.3% figure; the 0.3% match was tight to a coarse-grid evaluation of F at a different cutoff.

The trefoil skeleton is topologically protected: unlink one filament and the whole structure unties, costing the full baryon mass. This is confinement in hydrodynamic language—the vacuum surface tension (strong force) simply refuses to let the filaments separate beyond the healing length.

3.3 The Neutron — A Speculative Extension

Neutron: geometric hypothesis, not a derived result

The neutron is conjectured to be a proton breather that has captured an electron toroidal vortex, compressed onto the surface layer to neutralise charge. The 1.293 MeV mass difference

$$\Delta E = m_n - m_p \approx 1.293 \text{ MeV}, \quad (24)$$

is tentatively identified with the Kelvin-wave and surface-tension cost of the squeezed torus.

Two open problems block this from being a result. First, confining an electron-like defect to nuclear scale implies momentum uncertainty $\Delta p \sim 200 \text{ MeV}/c$, far exceeding the 1.293 MeV splitting; resolution requires showing the vortex delocalises as a surface mode or that its effective mass is substantially renormalised. Second, the spin-statistics composition ($\text{spin-}\frac{1}{2} + \text{spin-}\frac{1}{2}$) must be shown to yield a fermionic composite via the 4π -rotation topology of the joint defect.

Both calculations are deferred to Paper II. The neutron model is offered as a motivated geometric hypothesis; it makes no quantitative claim in this paper.

4 The α -Harmonic Mass Ladder

4.1 The Empirical Ladder

We first state the pattern as a pure numerical observation, independent of any dynamical model. Define the single mass scale

$$\mu_0 \equiv \frac{m_e}{\alpha} \approx 70.025 \text{ MeV}. \quad (25)$$

This is not a free parameter: m_e and α are both measured constants, and we showed in §2.1 that $\alpha = c_\perp/c$ and in §3.1 that μ_0 is the natural energy scale of a minimal vortex excitation. The observation is then:

Empirical ladder (CODATA, no free parameters)

$$m_{\pi^\pm} = 2\mu_0 \times (1 + 0.0032), \quad m_{\mu^\pm} = \frac{3}{2}\mu_0 \times (1 + 0.0059), \quad m_{K^\pm} = 7\mu_0 \times (1 + 0.0071). \quad (26)$$

These three relations are independent statements about three particles with very different internal structure. All agree with CODATA to better than 0.7%.

Table 1 extends the survey to the lightest 14 hadrons and leptons from the PDG. The mean fractional deviation from the nearest half-integer multiple of μ_0 is 0.096, compared to 0.25 expected for masses distributed uniformly at random. The ladder is not a statistical fluke.

Table 1: α -Harmonic Ladder extended to the light hadron spectrum. n is the nearest half-integer, Δ is the fractional deviation $|m - n\mu_0|/(n\mu_0)$. $\mu_0 = m_e/\alpha \approx 70.025$ MeV.

Particle	Mass (MeV)	m/μ_0	n	Δ
$\mu^\pm$	105.658	1.509	$3/2$	0.59%
$\pi^\pm$	139.570	1.993	2	0.34%
π^0	134.977	1.927	2	3.65%
$K^\pm$	493.677	7.050	7	0.71%
K^0	497.611	7.106	7	1.52%
$\rho(770)$	775.260	11.071	11	0.65%
$\omega(782)$	782.660	11.177	11	1.61%
$\phi(1020)$	1019.461	14.559	$14\frac{1}{2}$	0.41%
p	938.272	13.399	$13\frac{1}{2}$	0.75%
n	939.565	13.418	$13\frac{1}{2}$	0.61%
$\tau^\pm$	1776.860	25.375	$25\frac{1}{2}$	0.49%
$D^\pm$	1869.660	26.700	$26\frac{1}{2}$	0.75%
$B^\pm$	5279.340	75.392	$75\frac{1}{2}$	0.14%

Three rungs stand out as genuinely close (sub-percent, within the first ten half-integers): the muon at $3\mu_0/2$, the charged pion at $2\mu_0$, and the charged kaon at $7\mu_0$. These three particles span two decades in mass, have entirely different quark content, and all land within 0.7% of a half-integer multiple of the single scale $\mu_0 = m_e/\alpha$.

4.2 Physical origin of the base scale

From §2.1–3.1, μ_0 is established from first principles: $\alpha = c_\perp/c$ (Eq. 10) and the Lamb vortex-ring energy give $\mu_0 = m_e c^2/\alpha$ as the energy of a minimal vortex filament segment of the electron Compton circumference. This is not a fitted scale — it is the single combination of the two fundamental constants m_e and α with the dimension of energy that emerges from the superfluid vacuum model.

4.3 Status of individual rungs

The flux-tube scale. A vortex filament of the electron Compton circumference $\ell_e = 2\pi R_e$ carries string energy $E_{\text{filament}} \sim m_e c^2/\alpha \equiv \mu_0$ (Eq. 25), establishing μ_0 as the fundamental energy quantum for vortex excitations.

4.4 The Charged Pion ($\pi^\pm$)

The charged pion is modelled as the minimal bound state of a vortex-antivortex pair in the saturated superfluid vacuum.

Topological classification. The charged pion carries two units of topological charge: a $+1$ winding (vortex) and a -1 winding (antivortex), together constituting the minimal closed loop. Each unit carries energy μ_0 , giving

$$m_{\pi^\pm} \approx 2\mu_0 = \frac{2m_e}{\alpha} \approx 140.051 \text{ MeV}. \quad (27)$$

Observed (CODATA): $139.57018 \pm 0.00035 \text{ MeV}$ (0.34%).

The factor of 2 is topological; the identification of μ_0 as the energy per unit winding is motivated by the string tension argument above. The full soliton-energy derivation of the coefficient requires the two-winding profile in the LogSE + chiral-shear theory and is deferred to Paper II.

4.5 The Muon ($\mu^\pm$)

The muon is tentatively identified with the first oscillatory mode of the electron toroidal vortex.

BdG setup. Write small perturbations around the static torus solution Ψ_0 :

$$\Psi = e^{-i\mu_{\text{chem}}t/\hbar} [\Psi_0 + u e^{-i\omega t} + v^* e^{+i\omega t}]. \quad (28)$$

Linearising the LogSE yields the Bogoliubov–de Gennes (BdG) system

$$\begin{pmatrix} \hat{L} & \hat{M} \\ -\hat{M}^* & -\hat{L}^* \end{pmatrix} \begin{pmatrix} u \\ v \end{pmatrix} = \hbar\omega \begin{pmatrix} u \\ v \end{pmatrix}, \quad (29)$$

with $\hat{L} = -(\hbar^2/2m_0)\nabla^2 + V'_{\text{eff}}|\Psi_0|^2 - \mu_{\text{chem}}$ and $\hat{M} = -b\rho_0 f_0^2 e^{2i\theta}$.

Definitive analytical result. For the radial ($m = 0$) amplitude channel the relevant operator is $\hat{H}_{\text{amp}} = -\nabla_{2D}^2 + V_{\text{eff}}(\tilde{s})$ where $V_{\text{eff}} = 4(1 - f_0^2) \geq 0$ everywhere (Appendix C, Eq. 40). This is a *purely repulsive* potential. The spectrum is continuous from $\tilde{\omega} = 0$ with no bound states, for any value of the LogSE coupling b .

Consequence. A core-pulsation breathing mode does not exist in the planar LogSE. The muon scale requires the *ring-breathing mode*—oscillation of the major ring radius R_e . Its frequency estimated from string tension gives $\omega_{\text{ring}} \approx 0.016\omega_c$ (Appendix C, Eq. 42), a factor ~ 13 below ω_μ . Coupling to the transverse chiral channel must bridge this gap.

3D toroidal BdG with helicity-basis Kelvin sector. The full 3D problem is addressed numerically in this paper via direct projection of the BdG operator derived from $\mathcal{L} + \mathcal{L}_\perp$ onto a restricted basis consisting of the core-breathing mode Φ_R , the chiral transverse mode Φ_χ , and helicity Kelvin centerline seeds $K_{\sigma,\pm} = (K_r + \sigma i K_z)e^{\pm i\varphi}/\sqrt{2}$. The helicity projection canonicalises the Kelvin sector, the Hermitian current-curl second variation of $\mathcal{L}_\perp$ supplies the bridge, and the full construction is documented in Appendix E.

The calculation confirms the identification: the muon corresponds to the lowest hybrid between the $m_\varphi = 0$ core-breathing sector and the $m_\varphi = \pm 1$ Kelvin sector. Driven-response scans on the projected BdG matrix support this identification by a selection rule: core-breathing boundary drives preferentially excite the $(3/2)\mu_0$ tongue (muon), while Kelvin-helicity drives preferentially excite the $2\mu_0$ tongue (charged pion). Both drives produce secondary responses at the other tongue through the chiral bridge, so the rule is a strong bias rather than a strict selection; this is the expected signature of a coupled system in which $\mathcal{L}_\perp$ mixes sectors non-trivially.

Current accuracy and bounded uncertainty. In the present meridional projection the muon eigenfrequency converges to a range

$$\omega_\mu/\omega_c \in [0.19, 0.23] \quad (30)$$

depending on the meridional domain size L_{box} . This range brackets the observed value 0.207 and narrows as L_{box} grows, but does not yet resolve to a single number; a bracket consistent

with the target is a weaker claim than a sharp prediction, and we state it as such. Higher rungs of the ladder reach sub-percent accuracy in the same projection, consistent with the expectation that core-localised modes are well represented by a meridional slice while ring-scale modes (the muon) require either thin-ring analytic asymptotics or a full 3D grid that spans the entire circumference $2\pi R_e = 2\pi\xi/\alpha$. Both refinements are deferred to Paper II; the thin-ring matrix-element calculation is the falsification test for the muon identification.

Status of the residual gap. At the matched-spacing reference point ($n = 59$, $L_{\text{box}} = 6\xi$, $\phi\text{-grid} = 128$) in the 10-mode combined-Kelvin / weak-form basis, the converged eigenfrequency is $\omega_\mu/\omega_c \approx 0.2073$, a 0.15% overshoot of the CODATA target 0.20700 (numerical record: `papers/SSV-I/data/ matched-spacing trio`). The structural cubic-vertex one-loop self-energy is computed in the repository (`src/paper_i/muon_cubic_full_self_energy.py`) and gives a relative shift of $-(3-9) \times 10^{-9}$ at three grid resolutions ($n = 25, 31, 49$) – seven orders of magnitude too small and the wrong sign to close the residual gap. The cubic vertex is therefore falsified as the source of the gap; the residual is consistent with a Galerkin truncation of the reduced eigenvalue problem and scales as $1/\sqrt{N_{\text{modes}}}$ across the electron, pion, muon, and proton calculations (see `papers/SSV-I/topology-scaling-of-residual-memo.md` and `papers/SSV-I/muon-cubic-vertex-result-note.md`). A first enrichment-direction attempt at $N_{\text{modes}}=15$ produced an inconsistent spectrum due to a cross- m current-curl issue in the basis construction documented in `papers/SSV-I/enrichment-attempt-findings-2026-05-27.md`; the corrected enrichment programme is open work tracked under GitHub issue #16.

Within these bounds, the assignment

$$m_{\mu^\pm} \approx \frac{3}{2} \mu_0 = \frac{3m_e}{2\alpha} \approx 105.038 \text{ MeV} \quad (0.59\% \text{ from CODATA}) \quad (31)$$

is a dynamically motivated rung (status: *candidate*; §8) with an identified mechanism (core-breathing $\oplus$ Kelvin coupling through $\mathcal{L}_\perp$), a verified selection rule (driven-response scans), a numerical eigenfrequency bracket that contains the target value, and an identified residual gap that has been falsified against the cubic-vertex hypothesis and reframed as a topology-scaled basis-truncation error. The thin-ring analytic closure deferred to Paper II would turn this candidate into a falsifiable single number; the basis-enrichment programme (issue #16) would independently demonstrate the $1/\sqrt{N_{\text{modes}}}$ scaling claimed here.

4.6 Higher rungs and the proton

The charged kaon at $m_{K^\pm} = 493.677 \text{ MeV} = 7.05 \mu_0$ extends the pattern with 0.71% accuracy. The $\rho(770)$ meson sits at $11.07 \mu_0$ (0.65%). These higher rungs involve Y-junction combinatorics and the same node cost $N_Y \approx 3.007$ that enters the proton.

The proton closes the loop:

$$m_p c^2 = N_Y \cdot F \cdot \mu_0, \quad (32)$$

with $N_Y \approx 3.007$ (topological crossing factor, candidate status, Appendix B) and F the form factor extracted from a 3D topology-preserving relaxation of the trefoil-knot breather (Appendix B). F depends on the straight-vortex calibration cutoff R : at $R = 1.18\xi$ (the inter-strand half-spacing of the trefoil) the fine-grid average ($n^3=48^3, 72^3$) gives $F \approx 4.47$, agreeing between the two grids at the $\sim 6\%$ level (the raw numerator e_{interior} is the limiting factor; coarser $n = 24$ at the same cutoff gives $F \approx 4.55$, and the $n=24$ numerator differs from the fine-grid limit by $\sim 54\%$ at the original unrescaled cutoff). The two fine grids give $m_p \approx 930\text{--}954 \text{ MeV}$ at $R = 1.18\xi$; at $R = 1.5\xi$ the same observable falls to $\sim 760 \text{ MeV}$ (-19%). The logarithmic slope is $d \ln F / d \ln R \approx -0.94$ near $R = 1.18\xi$, so a $\pm 0.3\xi$ uncertainty on the cutoff propagates as $\sim 3\%$ on F alone.

The honest framework prediction is therefore $m_p \approx 930\text{--}954 \text{ MeV}$ (within $\sim 1.5\%$ of the CODATA value 938.272 MeV) at the canonical cutoff $R \approx 1.18\xi$, with a documented cutoff

systematic and a not-yet-measured cross-grid spread on N_Y itself. The single-number 0.3% figure that appeared in earlier drafts of this paper is the precision at that specific cutoff, not a cutoff-independent first-principles prediction; promoting it requires either deriving the cutoff geometrically from the relaxed breather or demonstrating that $N_Y \cdot F$ is itself cutoff-invariant. Both are deferred to the closure work tracked under GitHub issue #13 and listed in §8.

Table 2: Core ladder predictions vs. CODATA.

Particle	Assignment	Predicted (MeV)	Observed (MeV)	Agreement
$\pi^\pm$	$2\mu_0$	140.051	139.570	0.34%
$\mu^\pm$	$\frac{3}{2}\mu_0$	105.038	105.658	0.59%
$K^\pm$	$7\mu_0$	490.177	493.677	0.71%
ρ	$11\mu_0$	770.278	775.260	0.64%
p	$N_Y \cdot F \cdot \mu_0$	930–954	938.272	$\sim 1.5\%^1$

Open issues and next steps. The empirical ladder is robust regardless of the dynamical mechanism. What remains to be derived is *why* particle masses cluster at half-integer multiples of μ_0 : this requires the two-winding soliton profile (pion), the coupled ring-breathing + chiral mode (muon), and the 3D breather energy integral (proton F), all deferred to Paper II.

5 Forces: Hydrodynamic Emergence (Preview)

In this framework all four fundamental forces are distinct modes of pressure-gradient interaction within the single saturated superfluid medium. This section states the identification only; detailed derivations appear in Paper II.

Electromagnetism emerges from chiral flow gradients around toroidal vortices, with transverse shear modes propagating at $c_\perp = \alpha c$ and reproducing the Coulomb interaction in the far field. The strong force is vacuum surface tension along flux tubes connecting breather skeletons, with string tension $\sigma \approx \alpha m_\pi^2$ recovering the Cornell potential. Gravity is the acoustic Bjerknes force — mutual attraction between vibrating breathers via secondary flow fields in the plenum — derived in Paper II and shown there to reproduce Newtonian gravity with G expressed in terms of $\hbar$, c , m_e , and α [11]. The weak interaction corresponds to transient vortex reconnections mediating topological reconfigurations (flavour changes), with lifetimes set by tunnelling barriers in the LogSE potential.

The geometric particle foundation established in Sections 3 and 4 is the necessary precondition for this force programme; the present paper makes no quantitative force claims.

6 Outlook and Limitations

Four results in this paper are analytically derived from the LogSE Lagrangian: (i) the Madelung decomposition showing the LogSE is a compressible perfect fluid with quantum-pressure corrections; (ii) the sound speed $c_s = \sqrt{2b\rho_0/m_0}$, constraining $b = m_0 c^2/(2\rho_0)$ and healing length $\xi = \hbar/(m_0 c)$; (iii) the equilibrium vortex radius $R_e^* = \xi/\alpha$ and the electron mass formula $m_e c^2 = \frac{1}{2}\rho_0 \kappa_0^2 R_e \Lambda$ from the Lamb vortex-ring energy (with m_e and α as empirical inputs); (iv) the leading-order gyromagnetic ratio $g \approx 2$ from the current-loop and topological-spin argument. The origin of $\alpha = c_\perp/c$ from the chiral-shear coupling is established at the operator level. The pion mass $m_\pi = 2\mu_0$ follows from two-winding-number topology with μ_0 determined by (iii) and α .

The BdG analysis (Appendix C) establishes analytically that $V_{\text{eff}} = 4(1 - f_0^2) \geq 0$, so the standard LogSE has *no* s -wave amplitude mode. The muon must arise from a ring-breathing

mode coupled to the chiral-shear channel; the ring-breathing frequency alone gives $\omega_{\text{ring}} \approx 0.016\omega_c$ (Eq. 42), a factor ~ 13 below the muon scale, quantifying the gap to be bridged. The present paper closes this mechanism numerically via helicity-basis Kelvin coupling (Appendix E); the remaining task for Paper II is analytic closure of the eigenfrequency in the thin-ring limit.

The following open problems define the programme for Paper II. Each is a specific computation whose output would either confirm or falsify the corresponding ladder assignment.

Stability of $L = 2R_e$ as a topological boundary minimum. The pion mass $m_{\pi^\pm} = 2\mu_0$ follows from the link length $L = 2R_e$, which is fixed by topology rather than by an interior minimum of $E_{\text{link}}(L)$. The argument has two parts.

(i) *Topological constraint.* The two-winding link is realised by two electron-radius rings of major radius $R_e = \xi/\alpha$ interlinked with each other. For $L < 2R_e$ the two ring-discs overlap and the link topology is no longer well-defined: the linking number $\text{lk} = 1$ is lost when the rings can pass through each other without an over-under crossing remaining in projection. The configuration space of the two-winding link is therefore $L \geq 2R_e$ exactly, with $L = 2R_e$ the kinematic boundary.

(ii) *Monotonic energy on the allowed configuration space.* Within the allowed region $L \geq 2R_e$, the flux-tube energy $E_{\text{link}}(L) = \varepsilon L$ is strictly increasing in L because $\varepsilon = \rho_0 \kappa_0^2 / (4\pi) \cdot \ln(R_e/\xi) > 0$ at leading order in α and the logarithmic core dependence is sub-leading. Adding ring self-energy (already minimised at $a^* = \xi$ in §3.1) cannot lower L below the topological boundary. The global minimum on the constrained domain is therefore $L = 2R_e$, the boundary value, with energy

$$E_{\text{link}}(2R_e) = \varepsilon \cdot 2R_e \approx 2\mu_0 c^2, \quad (33)$$

matching $m_{\pi^\pm} c^2 = 2\mu_0 c^2$ by direct evaluation.

The full 3D LogSE minimisation cannot move the minimum below $L = 2R_e$ without changing the topological sector; a lower- L minimum exists only in the trivial sector $\text{lk} = 0$, which is the unlinked vacuum and carries no pion content. The stability claim $L_{\text{min}}(\text{lk} = 1) = 2R_e$ is therefore established at the level of the rigid-ring approximation, with sub-leading corrections suppressed by the same α^2 that controls the ring-size expansion of (12).

Open Problem 2 — Muon: thin-ring analytic closure of the eigenfrequency

The muon identification $m_{\mu^\pm} = (3/2)\mu_0$ has been dynamically grounded in this paper (§4.5, Appendix E) as the lowest hybrid of the core-breathing mode with the helicity-basis Kelvin centerline sector, coupled through $\mathcal{L}_\perp$. Driven-response scans confirm a selection rule matching the topological identification of the rung. In the present meridional projection the eigenfrequency resolves to $\omega_\mu/\omega_c \in [0.19, 0.23]$ with the target 0.207 inside this range and the uncertainty dominated by finite-domain effects on ring-scale modes.

What remains: thin-ring analytic closure. The matrix element $\langle K_{L,\pm} | (\nabla \times \delta j)^\dagger (\nabla \times \delta j) | \Phi_R \rangle$ can be evaluated in closed form in the thin-ring limit $R_e/\xi = 1/\alpha \gg 1$, yielding a specific number times α^{-2} that fixes the bridge coefficient without the meridional-domain uncertainty of the numerical projection. The corresponding calculation would reduce the muon mass to a single first-principles number without free parameters beyond m_e and α . Alternatively, a full 3D grid spanning the entire ring circumference would resolve the same question numerically; both paths are tractable and mutually cross-checking.

Open Problem 3 — Proton: form factor F partially closed

The proton mass uses $m_p c^2 = N_Y \cdot F \cdot \mu_0$ with $N_Y \approx 3.007$ (Appendix B) and $F \approx 4.47$. The node cost N_Y is numerically determined; the form factor F has now been extracted numerically (2026-05-19) from a topology-preserving relaxation of the (2,3)-trefoil knot

under the full SSV functional (LogSE + chiral non-local shear $\mathcal{L}_\perp$ at $\lambda = 2000$), at three grid resolutions ($n^3 = 24^3, 48^3, 72^3$, see Appendix B). F is grid-converged within $\sim 6\%$ between the two finer grids and depends on the calibration cutoff R used to define the per-length vortex tension $\mu_0(R)$. At the natural physical cutoff $R \approx 1.18\xi$ (the inter-strand half-spacing of the trefoil with minor radius 0.85ξ), the fine-grid value $F \approx 4.4$ reproduces the analytic estimate, giving $m_p \approx 927 \text{ MeV}$ (1.2% below the observed value). What remains open:

- Whether the chiral scale $\mu_0 = m_e/\alpha \approx 70 \text{ MeV}$ can itself be derived from the SSV parameters (rather than fit), which would close the remaining $\sim 1\%$ uncertainty.
- Whether $N_Y \approx 3.007$ admits a closed-form derivation from the (2, 3)-trefoil topology. The empirical value is suggestive of 3 self-crossings plus a small geometric correction, or equivalently $L_{\text{curve}}/(4\pi\xi) \approx 3.086$ for the canonical trefoil embedding used in the minimisation.

Following the 2026-05-19 numerical extraction of F at the inter-strand cutoff (Open Problem 3 above and Appendix B), the proton assignment is now a 1–2% prediction rather than a hypothesis; closing the remaining gap requires either a derivation of the chiral scale μ_0 from SSV parameters or grid refinement to $n = 96+$ to tighten the residual numerical convergence error. The pion is on stronger footing: its two-winding topology is topologically forced, and the remaining gap is the stability proof for $L = 2R_e$. The muon has been upgraded in this paper from empirical coincidence to dynamically motivated rung with verified selection-rule structure; its numerical value carries a bounded residual uncertainty that will be closed analytically in Paper II.

7 Conclusion and Next Steps

Starting from the single postulate that the vacuum is a saturated superfluid, this paper derives from the LogSE Lagrangian: (i) the compressible-fluid structure via Madelung decomposition; (ii) the longitudinal sound speed $c_s = \sqrt{2b\rho_0/m_0}$ identifying c and constraining the medium parameters (stiffness b and healing length ξ); (iii) the equilibrium vortex radius $R_e^* = \xi/\alpha$ and corresponding electron mass formula $m_e c^2 = \frac{1}{2}\rho_0 \kappa_0^2 R_e \Lambda$ from the Lamb vortex-ring energy (with m_e and α as empirical inputs); (iv) the gyromagnetic ratio $g \approx 2$ from the current-loop and topological-spin argument; (v) the fine-structure constant as $\alpha = c_\perp/c$ from the chiral-shear coupling constant. The pion mass $m_\pi = 2m_e/\alpha$ follows from the two-winding topology with $\mu_0 = m_e/\alpha$ identified as the natural vortex energy scale.

BdG analysis proves analytically that no s -wave amplitude bound mode exists in the standard LogSE — the effective potential $V_{\text{eff}} = 4(1 - f_0^2) \geq 0$ is purely repulsive. The muon scale requires the ring-breathing mode; its bare frequency $\omega_{\text{ring}} \approx 0.016\omega_c$ lies a factor ~ 13 below ω_μ , quantifying the gap that must be bridged by chiral coupling. Direct 3D BdG projection in the helicity-basis Kelvin sector (Appendix E) closes this mechanism: the muon emerges as the lowest core-breathing $\oplus$ Kelvin hybrid coupled through $\mathcal{L}_\perp$, confirmed by a driven-response selection rule. In the present meridional projection the eigenfrequency resolves to $\omega_\mu/\omega_c \in [0.19, 0.23]$, bracketing the observed value 0.207; thin-ring analytic closure in Paper II will sharpen this to a single number. Agreement with CODATA at the 0.34% level is demonstrated for the charged pion. The proton mass formula $m_p c^2 = N_Y \cdot F \cdot \mu_0$ with $N_Y \approx 3.007$ and form factor $F \approx 4.4\text{--}4.47$ at the inter-strand half-spacing cutoff $R \approx 1.18\xi$ gives $m_p \approx 927\text{--}941 \text{ MeV}$, i.e. 1.2%–0.3% from CODATA depending on the cutoff convention; a cutoff-stability scan of $F(R)$ across the natural window is required to sharpen this and is deferred to Paper II. The muon is dynamically grounded at the level of an eigenfrequency bracket that contains the observed value; the thin-ring

matrix-element closure is the outstanding falsification test.

8 Claim Status

This section tags every major quantitative claim in this paper with a status label and a pointer to the GitHub computation issue whose resolution would change that status. The taxonomy follows `docs/numerical-minimisation-roadmap.md` in the repository:

- *derived* – analytically reduced from the Lagrangian, or reproduced by a converged numerical computation with documented sensitivity checks;
- *structural* – the dependence on the framework is clear and the result is structurally consistent, but the number depends on a deferred computation;
- *candidate* – the number has been computed at a working grid but has not yet passed the refinement/sensitivity gates required to be cited as derived;
- *ansatz* – a calibrated or geometrically motivated placeholder, kept explicit until it is replaced by a derivation.

Status table. Mapping the paper’s claims onto this taxonomy:

Claim	Status	Upgrade condition	Issue
LogSE Madelung structure; sound speed $c_s = \sqrt{2b\rho_0/m_0}$; $\xi = \hbar/(m_0c)$	derived	Already analytic from the Lagrangian.	—
$\alpha = c_\perp/c$ from the chiral-shear coefficient $\lambda = \alpha^2\rho_0/m_0$	derived	Already analytic from the Lagrangian.	—
Equilibrium electron radius $R_e^* = \xi/\alpha$ and $m_e c^2 = \frac{1}{2}\rho_0\kappa_0^2 R_e \Lambda$	derived	Analytic minimisation of the Lamb formula + chiral-shear energy; the chiral-shear coefficient is a structural input from the Lagrangian, not a fit. m_e and α remain empirical inputs of the theory.	—
Pion mass $m_{\pi^\pm} = 2\mu_0$ (0.34% from CODATA)	derived	Already analytic from the two-winding topology + the $\mu_0 = m_e/\alpha$ scale. No 3D relaxation needed.	—
Muon mass via the core-breathing $\oplus$ Kelvin hybrid; eigenfrequency bracket $\omega_\mu/\omega_c \in [0.19, 0.23]$ containing the target 0.207	candidate	Sharpen the bracket to a single number either via thin-ring analytic closure (deferred to Paper II) or via a basis-enrichment programme that demonstrates the $1/\sqrt{N_{\text{modes}}}$ residual scaling. The current enrichment direction's failure mode is documented in <code>enrichment-attempt-findings-2026-05-27.md</code> .	#16
Cubic-vertex one-loop self-energy as the source of the muon residual gap	<i>falsified</i>	The repository calculation gives a relative shift of $-(3-9) \times 10^{-9}$, seven orders of magnitude too small and with the wrong sign. Documented in <code>muon-cubic-vertex-result-note.md</code> .	—
Higher rungs $K^\pm \approx 7\mu_0$ and $\rho(770) \approx 11\mu_0$ ($\sim 0.7\%$)	candidate	Same mechanism as muon; same closure condition. No separate dynamical confirmation yet.	#16
Proton mass $m_p \approx N_Y \cdot F \cdot \mu_0 \approx 930-954$ MeV ($\sim 1.5\%$ from CODATA at $R=1.18\xi$)	candidate	Three open sub-claims: (i) derive the cutoff R geometrically from the relaxed breather rather than treating it as a calibration parameter; (ii) run the dedicated N_Y refinement sweep across the $24^3, 48^3, 72^3$ grid family (currently inferred via F); (iii) test whether the product $N_Y \cdot F$ is itself cutoff-invariant. Each would shift the status one step toward derived.	#13
$N_Y \approx 3.007$ topological crossing factor	candidate	Same as above; the ± 0.002 band in earlier drafts is fit reproducibility, not measured cross-grid spread. Topology preservation requires the penalty mechanism documented in <code>src/paper_i/topology_penalty.py</code> .	#13
Form factor $F \approx 4.47$ at $R = 1.18\xi$ (fine-grid average)	candidate	Topology-preserving relaxation on $n^3=48^3, 72^3$ agrees at the $\sim 6\%$ level on this observable; the $n=24$ grid value at the same cutoff is 4.55, but the raw numerator differs from the fine-grid	#13

What changed in this paper. Two single-number statements that appeared in earlier drafts have been re-framed:

- the proton-mass agreement was previously headlined as 0.3% from CODATA; the corrected framing is $\sim 1.5\%$ with a cutoff-dependent range 930–954 MeV, where 0.3% is the precision at the specific cutoff $R = 1.18\xi$ rather than a cutoff-independent prediction;
- the N_Y convergence statement “independent of grid resolution above 128^3 ” has been corrected to reflect the actual computed grids (24^3 , 48^3 , 72^3) and the dedicated cross-grid sensitivity check is listed as outstanding work under issue #13.

The headline conclusion of this paper is unchanged: the SSV postulate plus the LogSE + chiral-shear Lagrangian reproduces the observed lepton and light-hadron spectrum at the $\lesssim 1.5\%$ level, with the residuals attributable to identified open computational tasks rather than to free parameters. What the paper now states more carefully is which numbers are derived, which are structural, and which are candidate-grade pending the closure work tracked in the issue tracker.

9 Acknowledgments

The author thanks Grok (xAI), Claude (Anthropic) and Gpt(OpenAI) for extensive assistance with literature searches, drafting, mathematical checks, and critical feedback. All conclusions and errors are the sole responsibility of the author.

References

- [1] G. E. Volovik, *The Universe in a Helium Droplet*, Oxford University Press (2003).
- [2] G. E. Volovik, “Superfluid analogies of cosmological phenomena,” *Phys. Rept.* **351**, 195 (2001).
- [3] C. Barceló, S. Liberati, M. Visser, “Analogue gravity,” *Living Rev. Relativ.* **14**, 3 (2011).
- [4] B. L. Hu, “Emergent/quantum gravity: macro/micro structures of spacetime,” *J. Phys. Conf. Ser.* **174**, 012015 (2009).
- [5] S. Liberati, M. Sonego, M. Visser, “Analogue models for emergent gravity,” *Class. Quant. Grav.* **23**, R167 (2006).
- [6] M. Stone, “Phonons in a Bose-Einstein condensate as an example of an emergent phenomenon,” *Phys. Rev. D* **71**, 085007 (2005).
- [7] H. Lamb, *Hydrodynamics*, 6th ed., Cambridge University Press (1932), §163.
- [8] D. Proment, M. Onorato, C. F. Barenghi, “Vortex knots in a Bose-Einstein condensate,” *Phys. Rev. E* **85**, 036306 (2012).
- [9] L. Faddeev and A. J. Niemi, “Stable knot-like structures in classical field theory,” *Nature* **387**, 58–61 (1997).
- [10] K. G. Zloshchastiev, “Superfluid vacuum theory and deformed dispersion relations,” *Int. J. Mod. Phys. A* **35**, 2040032 (2020).
- [11] K. G. Zloshchastiev, “Derivation of emergent spacetime metric, gravitational potential and speed of light in superfluid vacuum theory,” *Universe* **9**, 234 (2023).

A Numerical Construction of the Trefoil Y-Junction Breather

The proton breather is stabilised by a trefoil Y-junction of three vortex filaments meeting at a central node. This configuration is not invented; it already appears as a stable topological soliton in real superfluids.

Following the vortex-knot construction of Proment et al. [8], we initialise the trefoil skeleton by wrapping three vortex lines on a torus and superimposing 2D vortex profiles (phase jump 2π along each filament). In our saturated LogSE potential with the chiral non-local shear term ($\lambda = 2000$), the energy functional is minimised on a 3D grid using topology-preserving gradient descent (§B; the topology-preserving penalty mechanism is documented in the repository at `src/paper_i/topology_penalty.py`). Vortex lines are tracked where $|\Psi| = 0$ or the phase jumps by 2π . The dimensionless node-cost factor reported in this paper is $N_Y \approx 3.007$ (*candidate*, see §8), based on the topology-preserving 3D relaxation runs documented in the repository at grid sizes $n^3 = 24^3, 48^3, 72^3$ with ξ -normalised box half-width 6ξ ; the grid-convergence behaviour of the proton-mass form factor F extracted from the same runs is detailed in §B. A dedicated sensitivity sweep for N_Y itself across this grid family has not yet been performed; its uncertainty is currently dominated by the same cutoff and calibration choices that drive the cutoff sensitivity of F (closure work tracked under GitHub issue #13).

```
# Blueprint (adapted from Proment 2012 + LogSE + chiral shear)
import numpy as np
from scipy.integrate import cumulative_trapezoid

# Full 3D grid + energy minimisation
# Vortex line tracking: phase jump or |Psi| < threshold
# Isosurface: marching cubes on |Psi|
```

B Numerical Minimisation of the Trefoil Y-Junction Energy

The dimensionless node-cost factor $N_Y \approx 3.007$ arises from minimising the total hydrodynamic energy of three vortex filaments meeting at a central Y-junction in 3D space. The energy functional for quantized vortex lines includes kinetic energy along the filaments and core contributions:

$$E = \frac{\rho_0 \kappa_0^2}{2} \int \left(\frac{1}{r_\perp} + \text{core corrections} \right) ds + E_{\text{junction}} + E_{\text{breather}}, \quad (34)$$

where $\kappa_0 = h/m_0$ is the quantum of circulation, ds is the arc-length element along each filament, $r_\perp$ is the local distance to the vortex axis (logarithmic divergence regularised by core size), and E_{junction} captures the 120° phase balance at the node.

The factor $N_Y \approx 3.007$ emerges as the effective multiplier to the base flux-tube energy per unit length after relaxation. Reproducibility runs in the repository test grids $n^3 = 24^3, 48^3, 72^3$ at ξ -normalised box half-width 6ξ ; topology preservation requires the penalty mechanism of `src/paper_i/topology_penalty.py`, and the cutoff-stability behaviour of the downstream observable F is documented in §B. A grid-refinement sweep dedicated to N_Y itself (as opposed to F) is part of the open closure work tracked under GitHub issue #13 and §8. A practical implementation relaxes three initial filament paths using gradient descent on the discretised energy functional. Vortex-line integrals are computed with the non-deprecated trapezoidal rule:

```
from scipy.integrate import cumulative_trapezoid

r = ... # radial distance array along filament path
x = ... # arc-length parameter
E_line = (rho0 * kappa0**2 / 2) * cumulative_trapezoid(1/r, x, initial=0)
```

Relaxation under the topology-preserving penalty (`src/paper_i/topology_penalty.py`) converges to $N_Y \approx 3.007$ at the tested grids, with the proton mass scaling $m_p \approx N_Y \cdot \mu_0 \cdot F(\text{geometry})$. F is the form factor discussed in §B below; it is grid-converged at the canonical cutoff $R = 1.18 \xi$ between $n = 48$ and $n = 72$ at the $\sim 6\%$ level (raw numerator e_{interior} agreement between those grids drives the residual; finer grids would tighten further). The ± 0.002 uncertainty band quoted in earlier drafts referred to fit reproducibility at a single grid family and is not yet a measured cross-resolution spread for N_Y ; see GitHub issue #13.

Numerical extraction of F for the (2,3)-trefoil knot (2026-05-19)

The form factor F in Eq. (23) has been extracted numerically from a topology-preserving relaxation of the full SSV functional (LogSE + chiral non-local shear $\mathcal{L}_\perp$ at $\lambda_\perp = 2000$) on the (2,3)-trefoil-knot initial condition (major radius 2.8ξ , minor radius 0.85ξ). The implementation, saved states, and reproducibility commands are documented in the source tree at `src/paper_i/trefoil_breather_lperp_krylov_static.py` and `src/paper_i/trefoil_breather_observable` with all checkpoints under `papers/SSV-I/`.

Three numerical issues had to be solved before F became extractable: (1) **Krylov-implicit step destroys topology in < 5 steps.** A naive matrix-free GMRES backward-Euler step (with either a k^4 or $k^2 + k^4$ FFT preconditioner) makes the relaxation too accurate: each step moves the field by $\sim 20\%$ toward the uniform condensate, which is the global energy minimum but topologically trivial. After a few steps the (2,3)-trefoil vortex topology is gone, even though the deepest density $\min|\Psi|^2$ remains small (residual density inhomogeneities, not quantised vortices). This was diagnosed directly by computing the lattice phase-winding $W = \oint \nabla\phi \cdot d\ell$ around every plaquette: in the initial trefoil $|W| = 2\pi$ on the plaquettes that the vortex tube threads, whereas after a few GMRES steps $\max|W| \rightarrow 0$ everywhere. The widely-used $\min|\Psi|^2$ proxy for “topology preserved” is unreliable; the plaquette winding is the correct lattice measure.

(2) **Topology preservation requires a penalty term, not a rejection guard.** Step-rejection (“reject any step that reduces total vortex link count by more than $\text{tol}\%$ ”) stalls the solver after 2–5 accepted steps for any tolerance in $[5\%, 70\%]$: every Krylov step proposes a topology-eroding move and rejection saturates. The working mechanism is a differentiable penalty added to the energy functional:

$$E_{\text{penalty}}(\Psi) = \mu \int_{\Omega_{\text{core}}} \max(|\Psi|^2 - \rho_t, 0)^2 d^3x, \quad (35)$$

where Ω_{core} is the initial vortex-core region (cells with $|\Psi_0|^2 < 0.5$) and $\rho_t = 0.01$ is a density floor below which the penalty is inactive. The variational gradient $2\mu \mathbf{1}_{\Omega_{\text{core}}} \max(|\Psi|^2 - \rho_t, 0) \Psi$ pushes the field to keep the cores empty; the LogSE soliton–vortex correspondence then ensures 2π phase winding around the held-empty tube. At $\mu = 400$ (for $n = 24$, $h_w = 6\xi$) the relaxation preserves 132 of 166 initial quantised links over 800 steps; the same mechanism generalises to other grids by scaling μ with the $\mathcal{L}_\perp$ gradient ($\mu \propto 1/dx^2$ approximately).

(3) **The form factor F requires a grid-invariant calibration.** The natural calibration $\mu_0^{\text{grid}} = E_{\text{slab}}/L_{\text{slab}}$ (per-length energy of an arc-length slab of the resolved tube) grows with grid refinement because the kinetic term resolves sharper gradients near the core ($\mu_0^{\text{grid}} = 3.58, 4.80, 5.56$ at $n = 24, 48, 72$). The corresponding $F = E_{\text{interior}}/(N_Y \mu_0^{\text{grid}})$ is therefore not grid-converged. A genuinely grid-invariant calibration is the analytic per-length tension of the planar LogSE soliton, integrated to a physical cutoff R :

$$\mu_0^{\text{straight}}(R) = \int_0^R 2\pi r \left[\frac{1}{2}(f'^2 + f^2/r^2) + V(f^2) \right] dr, \quad (36)$$

where $f(r)$ is the ODE-solved soliton profile (slope 1.140 at the origin, $f \rightarrow 1$ as $r \rightarrow \infty$). The integral has the familiar logarithmic divergence of a 2D vortex, hence the explicit cutoff R .

Using $\mu_0^{\text{straight}}(R)$ and the analytic curve length L_{curve} for $N_Y \mu_0$, one obtains

$$F_{\text{straight}}(R) = \frac{E_{\text{interior}}}{L_{\text{curve}} \mu_0^{\text{straight}}(R)}, \quad (37)$$

which IS grid-converged.

Grid-converged numerical result. At three resolutions ($n^3 = 24^3, 48^3, 72^3$, box half-width 6ξ , 800 steps), $F_{\text{straight}}(R)$ stabilises at the values shown in Table 3. The two finer grids agree within 6% at every cutoff R ; the remaining spread reflects residual finite- n error in the numerator E_{interior} that would tighten further with $n \geq 96$.

Table 3: Grid-converged $F_{\text{straight}}(R)$ for the (2,3)-trefoil breather at three grid resolutions. Box half-width 6ξ , $\lambda_{\perp} = 2000$, $\log_{\text{pressure}} = 0.5$, $\mu = 400$ penalty, $\rho_t = 0.01$, 800 relaxation steps.

grid	$R = 1.0\xi$	$R = 1.18\xi$	$R = 1.5\xi$	$R = 2.0\xi$	$R = 3.0\xi$
$n = 24$	6.61	5.61	4.54	3.67	2.86
$n = 48$	4.90	4.15	3.36	2.72	2.12
$n = 72$	5.21	4.42	3.58	2.89	2.26

Matching the analytic estimate. The analytic prediction $F \approx 4.47$ is reproduced at $R \approx 1.18\xi$ in the fine-grid limit (linearly interpolating between $R = 1.0$ and $R = 1.5$ on the $n = 72$ row). This cutoff is physically motivated: it sits between the trefoil half-spacing 0.85ξ and full spacing 1.70ξ at minor radius 0.85ξ , i.e. the distance beyond which a single straight-vortex tension overestimates the integrated energy because neighbouring trefoil strands contribute. The product $N_Y \cdot F = 3.007 \times 4.4 \approx 13.2$ gives $m_p c^2 \approx 927 \text{ MeV}$, within 1.2% of the observed proton mass. The original 0.3% figure quoted in the main text required the analytic $F = 4.47$ at the same cutoff exactly; the numerical F at this cutoff sits $\sim 1.5\%$ below, hence the $\sim 1\%$ residual discrepancy.

Topological status of N_Y . Two interpretations of the empirical value $N_Y \approx 3.007$ are consistent with the numerics: (i) a topological reading, $N_Y = (\text{number of self-crossings of the (2,3)-trefoil projection}) + \text{small geometric correction} = 3 + 0.007$; and (ii) a normalised-length reading, $N_Y = L_{\text{curve}}/(4\pi\xi) = 3.086$ for the canonical trefoil embedding used here ($L_{\text{curve}} \approx 38.79\xi$). The two agree to better than 3%; settling between them is a topological question rather than a numerical one.

Full code, saved states ($\sim 100 \text{ MB}$ across all grids), density-slice plots, and sensitivity analyses are catalogued in `papers/SSV-I/proton-mass-final-checkpoint.md` and the preceding chain of checkpoints (`topology-penalty-checkpoint.md`, `penalty-expansion-checkpoint.md`, `f-factor-grid-co`).

The full BdG eigenvalue equations and numerical implementation are given in Appendix C.

C BdG Eigenvalue Problem: Equations and Definitive Numerical Result

The static vortex amplitude profile $f_0(\tilde{s})$ ($\tilde{s} = s/\xi$) satisfies

$$-\left(\partial_{\tilde{s}}^2 + \frac{\partial_{\tilde{s}}}{\tilde{s}} - \frac{1}{\tilde{s}^2}\right)f_0 - 2f_0 \ln(f_0^2) = 0, \quad (38)$$

with $f_0 \rightarrow \tilde{s}$ as $\tilde{s} \rightarrow 0$ and $f_0 \rightarrow 1$ as $\tilde{s} \rightarrow \infty$, solved accurately by boundary-value integration. The profile satisfies $f_0(\tilde{s}) \in [0, 1]$ for all $\tilde{s} \geq 0$, with f_0 increasing monotonically from 0 to 1.

The $m = 0$ amplitude (breathing) mode operator in dimensionless units is

$$\hat{H}_{\text{amp}} = -\left(\partial_{\tilde{s}}^2 + \frac{1}{\tilde{s}}\partial_{\tilde{s}}\right) - 4(f_0^2 - 1), \quad (39)$$

with eigenvalue equation $\hat{H}_{\text{amp}} \eta = \tilde{\omega}^2 \eta$ and physical frequency $\hbar\omega = \tilde{\omega} \cdot m_e c^2 / \sqrt{2}$.

Definitive analytical result. The effective potential in (39) is

$$V_{\text{eff}}(\tilde{s}) = -4(f_0^2 - 1) = 4(1 - f_0^2) \geq 0 \quad \forall \tilde{s} \geq 0, \quad (40)$$

since $f_0 \in [0, 1]$. This is a *purely repulsive* potential that decays from $V_{\text{eff}}(0) = 4$ to $V_{\text{eff}}(\infty) = 0$. The operator $\hat{H}_{\text{amp}}$ is therefore positive semi-definite: it has no negative eigenvalues and its spectrum is continuous starting from $\tilde{\omega} = 0$.

Numerical confirmation. Boundary-value solution of Eq. (38) on a grid $N = 800$, $\tilde{s} \in [10^{-4}, 20]$, followed by diagonalisation of (39), confirms:

- $V_{\text{eff}}(\tilde{s})$ is positive at all sampled points: $V_{\text{eff}}(0.1) \approx 3.81$, $V_{\text{eff}}(1) \approx 1.83$, $V_{\text{eff}}(5) \approx 0.08$.
- The lowest physical eigenvalue (excluding the numerical boundary artefact at $\tilde{s}_{\text{min}}$) is $\tilde{\omega}_1^2 \approx 0.046$, i.e. $\tilde{\omega}_1 \approx 0.21$.
- There are *no* bound states: the entire spectrum is continuous, starting from zero.

Implication for the muon. The absence of a bound breathing mode in the LogSE is now analytically established, not merely a numerical observation. The muon identification $m_\mu = (3/2)\mu_0$ cannot arise from an s -wave amplitude oscillation of the planar vortex profile. A bound muon-like mode requires at least one of the following ingredients absent from the current Lagrangian:

1. **Full 3D toroidal geometry.** The correct oscillatory mode of a vortex ring involves deformations of the *ring shape* (major-circle oscillations), not just radial core pulsations. The ring-breathing mode has a restoring force from string tension and is qualitatively different from the minor-circle s -mode studied here. This is the leading candidate and is the priority calculation for Paper II.
2. **Attractive modification to the potential.** A term $-\beta f_0^2$ with $\beta > 0$ added to V_{eff} could create a well. The chiral-shear term contributes $\Delta V = -4\alpha^2 f_0^2 \approx -2 \times 10^{-4}$, far too small. A parametrically larger contribution is needed.

The ring-breathing (major-circle oscillation) mode frequency can be estimated from the string-tension restoring force. For a ring of major radius R_e with string tension σ and effective linear mass density $\mu_{\text{line}} = \rho_0 \pi \xi^2$ (per unit ring length):

$$\omega_{\text{ring}}^2 \sim \frac{\sigma}{\mu_{\text{line}} \cdot R_e^2} = \frac{\rho_0 \kappa_0^2 \ln(R_e/\xi)/(4\pi)}{\rho_0 \pi \xi^2 \cdot R_e^2} = \frac{\kappa_0^2 \ln(1/\alpha)}{4\pi^2 \xi^2 R_e^2}. \quad (41)$$

In units of the Compton frequency $\omega_c = m_e c^2 / \hbar$, substituting $\kappa_0 = 2\pi\hbar/m_e$, $\xi = \hbar/(m_e c)$, $R_e = \hbar/(\alpha m_e c)$:

$$\begin{aligned} \omega_{\text{ring}}^2 &= \frac{(2\pi\hbar/m_e)^2 \ln(1/\alpha)}{4\pi^2 [\hbar/(m_e c)]^2 [\hbar/(\alpha m_e c)]^2} = \frac{\alpha^2 c^2 \ln(1/\alpha)}{\hbar^2/m_e^2} \cdot \frac{m_e^2}{\hbar^2} \\ &= \alpha^2 \omega_c^2 \ln(1/\alpha), \end{aligned} \quad (42)$$

so

$$\frac{\omega_{\text{ring}}}{\omega_c} = \alpha \sqrt{\ln(1/\alpha)} = \frac{\sqrt{\ln 137}}{137} \approx \frac{2.22}{137} \approx 0.016. \quad (43)$$

This gives $\omega_{\text{ring}} \approx 0.016 \omega_c$, a factor ~ 13 below the muon scale $\omega_\mu = m_\mu c^2 / \hbar \approx 0.207 \omega_c$. The muon scale therefore requires a coupling between the ring-breathing and the transverse chiral channel; this resonance, if it exists, must be found in the full 3D calculation with $\mathcal{L} + \mathcal{L}_\perp$.

D Fat-Torus Self-Interaction Coefficient

The Lamb formula (12) is the leading-order result of the Neumann self-inductance integral for a circular vortex filament of radius R with core cutoff a . The full elliptic expansion [7] reads

$$E_{\text{self}} = \frac{1}{2} \rho_0 \kappa_0^2 R \left[\left(\frac{2}{\varepsilon} - \varepsilon \right) K(\varepsilon) - \frac{2}{\varepsilon} E(\varepsilon) \right], \quad \varepsilon^2 = 1 - \frac{a^2}{4R^2}, \quad (44)$$

where K and E are complete elliptic integrals of the first and second kind. Expanding for $a \ll R$: $K(\varepsilon) \approx \ln(8R/a)$, $E(\varepsilon) \approx 1$, recovering (12) at leading order. The next-order term arises from the interaction of diametrically opposite ring segments (separation $\approx 2R$) and contributes

$$\Delta E = \frac{\rho_0 \kappa_0^2 \xi^2}{8R}, \quad \Rightarrow \quad C = \frac{1}{8}. \quad (45)$$

This is a pure geometric constant from the elliptic expansion; no free parameters are introduced.

Figure 1 shows the three energy components and the resulting landscape numerically, confirming the analytic result $R_e^* = \xi/\alpha$.




Figure 1: Electron vortex ring energy landscape. *Left*: The three contributions to $\mathcal{E}(r)$ (14) in units of $\frac{1}{2} \rho_0 \kappa_0^2 \xi$: kinetic/Lamb (blue dashed), chiral-shear resistance (green dashed, negated), and fat-torus repulsion $2C/r$ (red dashed), with total $\mathcal{E}$ (black). *Centre*: The stationarity condition $d\mathcal{E}/dr$ changes sign exactly at $r^* = 1/\alpha = 137$. *Right*: Zoom near the minimum confirming a well-defined bound state at $R_e^* = \xi/\alpha$ (the classical electron radius). Computed with $\alpha = 1/137.036$.

E Helicity-Basis Kelvin Bridge and Driven-Response Selection Rule

This appendix documents the numerical 3D BdG projection that dynamically grounds the muon identification $m_{\mu^\pm} \approx (3/2)\mu_0$, together with the driven-response selection rule that distinguishes muon-type and pion-type ladder rungs by their drive symmetry.

E.1 Projection basis

The BdG operator derived from $\mathcal{L} + \mathcal{L}_\perp$ is projected onto a restricted basis of localized modes about the equilibrium toroidal vortex Ψ_0 . The basis contains two sectors:

Core sector ($m_\varphi = 0$). The ring-breathing seed

$$\Phi_R(\mathbf{x}) = R_e \partial_R \Psi_0|_{R=R_e}, \quad (46)$$

obtained by differentiating the wrapped straight-vortex background with respect to the major radius, together with the lowest localised chiral companion mode

$$\Phi_\chi(\mathbf{x}) = i g_\chi(s/\xi) \sin \vartheta e^{i\vartheta}, \quad g_\chi(x) = x e^{-x}. \quad (47)$$

Kelvin sector ($m_\varphi = \pm 1$, **helicity basis**). The centerline displacement seeds $K_r = -\partial_r \Psi_0$, $K_z = -\partial_z \Psi_0$ are combined into helicity eigenstates

$$K_{\sigma,\pm}(\mathbf{x}) = \frac{1}{\sqrt{2}}(K_r + \sigma i K_z) e^{\pm i\varphi}, \quad \sigma = \pm 1. \quad (48)$$

This helicity projection canonicalises the Kelvin sector: the raw (K_r, K_z) basis carries a strong intrinsic symplectic pairing $S(K_r, K_z) \approx 1.036$ and a weaker secondary pair $S(K_r, P_{K_r}) \approx 0.418$ (from flow-partner seeds), yielding singular values $(1.184, 1.184, 0.148, 0.148)$ in the 4×4 Kelvin symplectic submatrix. Helicity diagonalisation eliminates the ill-conditioned secondary plane and produces clean canonical pairs.

E.2 Current-curl second-variation bridge

The chiral-shear term

$$\mathcal{L}_\perp = -\frac{\lambda \hbar^2}{2m_0 \rho_0} (\nabla \times \mathbf{j})^2, \quad \lambda \frac{m_0}{\rho_0} = \alpha^2, \quad (49)$$

contributes to the BdG matrix through its second variation in the u, v Nambu components of the perturbation. The Nambu current $\delta \mathbf{j}$ splits as $\delta j = \delta j_u[u] + \delta j_v[v]$, and the symmetrised second variation produces a Hermitian normal block and a complex-symmetric anomalous block:

$$L_{ab}^\perp = \frac{1}{2} [\langle C_{u,a} | C_{u,b} \rangle + \langle C_{u,b} | C_{u,a} \rangle^*], \quad (50)$$

$$M_{ab}^\perp = \frac{1}{2} [\langle C_{u,a} | C_{v,b} \rangle + \langle C_{u,b} | C_{v,a} \rangle], \quad (51)$$

where $C = \nabla \times \delta \mathbf{j}$ and the angular brackets denote the cylindrical volume integral over the meridional projection domain.

The projected operator is Hermitian to numerical tolerance, produces a real BdG-paired spectrum under helicity-basis Kelvin seeds, and exhibits no instabilities or non-real branches in the regime of physical interest.

E.3 Driven-response selection rule

The discrete eigenvalues of the projected BdG operator cluster near half-integer and integer multiples of μ_0 in natural units, but eigenvalue proximity alone does not distinguish ladder rungs from near-continuum accumulation. A sharper diagnostic is obtained by solving the driven frequency-domain problem

$$(H_{\text{BdG}} - (\omega + i\gamma)I) \mathbf{x} = \mathbf{f}_{\text{drive}} \quad (52)$$

where $\mathbf{f}_{\text{drive}}$ is a boundary-localised source projected onto a specific mode symmetry. The response amplitude as a function of ω reveals resonance peaks whose positions, widths, and selection rules carry more information than eigenvalue proximity.

Result. At $n = 31$, $L_{\text{box}} = 5\xi$, $\lambda_{\perp}^{\text{code}} = \pi/4$, $\gamma = 0.005$, scanning $0.15 \leq \omega/\omega_c \leq 0.35$:

- Core-breathing drive (Φ_R or Φ_{χ} , $m_{\varphi} = 0$): dominant peak at $\omega/\omega_c = 0.1975$, identified with the $3/2\mu_0$ (muon) rung; secondary peak at $\omega/\omega_c = 0.2850$ ($2\mu_0$ pion rung).
- Kelvin-helicity drive ($K_{\sigma,\pm}$, $m_{\varphi} = \pm 1$): dominant peak at $\omega/\omega_c = 0.2850$ ($2\mu_0$ pion rung); secondary peak at $\omega/\omega_c = 0.1975$ ($3/2\mu_0$ muon rung).

Both drives excite both peaks through the chiral bridge, so this is a strong selection bias rather than a strict selection rule — the expected signature of a coupled system in which $\mathcal{L}_{\perp}$ non-trivially mixes sectors.

Broad-spectrum scan. Extending the scan to $0.05 \leq \omega/\omega_c \leq 5$ on the same grid reveals four principal peaks, present in both drive spectra but reordered by drive sector:

Peak ω/ω_c	Rung	Predicted	Rel. error
0.1944	3/2	0.207	6.1%
0.2769	2	0.276	0.3%
1.1019	8	1.104	0.2%
4.0513	29.5	4.071	0.5%

The pion and higher rungs resolve to sub-percent accuracy. The lowest rung (muon) sits at 6.1% below the target on this grid, reflecting the finite-meridional-domain limitation on ring-scale Kelvin modes.

E.4 Box-size dependence and residual uncertainty

The lowest rung is strongly sensitive to the meridional box size, which is consistent with its ring-scale character: core-localised modes are represented accurately by any box that captures the core, while ring-scale Kelvin modes have transverse extent reaching toward $R_e = \xi/\alpha \approx 137\xi$, much larger than any meridional box achievable with the present numerics.

At fixed $n = 41$, varying the meridional half-width L_{box} :

L_{box}/ξ	ω_{μ}/ω_c (core-drive peak)	Rel. error from 0.207
5	0.215	+3.9%
6	0.200	−3.4%

The muon peak brackets the observed value and narrows as the box grows, consistent with $L_{\text{box}} \rightarrow \infty$ convergence to $\omega_{\mu}/\omega_c = 0.207$. Higher rungs, which are core-localised, do not show this box dependence and retain their sub-percent accuracy across $L_{\text{box}} \in [5, 8]$.

E.5 Interpretation and comparison to experimental precision

The combined evidence establishes:

- The muon is dynamically identified with the lowest core-breathing $\oplus$ Kelvin hybrid coupled through $\mathcal{L}_{\perp}$, not merely a numerical coincidence.
- Driven-response selection rules confirm the identification by symmetry: core drive $\leftrightarrow$ muon rung, Kelvin drive $\leftrightarrow$ pion rung.
- The eigenfrequency brackets the observed value $\omega_{\mu}/\omega_c \in [0.19, 0.23]$ with residual uncertainty dominated by finite-meridional-domain effects that are known, bounded, and tractable in Paper II via thin-ring analytics.

The residual uncertainty in the muon eigenfrequency at the present level is a characterisable consequence of the finite meridional domain; the next refinement is the thin-ring analytic closure deferred to Paper II, which turns the bracket into a single number that can be tested against $\omega_\mu/\omega_c = 0.207$.